

Outpainting Audio by Reusing What the Model Already Knows: A Minimal Full-Finetune Recipe for Latent Rectified-Flow Audio Diffusion

Francesco Colonnese
github.com/mrcolo/stems-loop

Abstract

Audio outpainting — continuing a recording past its end — is usually framed as a new capability requiring new machinery: extra input channels, adapter branches, or a separately trained continuation head. We observe that a modern latent audio diffusion model already contains everything the task needs. Stable Audio 3 Medium is trained with inpainting conditioning, taking a binary mask and a masked latent as local-additive conditions, and its mask distribution already includes a causal mode that retains a random-length prefix. Outpainting is that mode. Consequently a finetune for outpainting can change no architecture at all: no new channels, no adapters, no surgery on the input projection. Only the transformer weights move.

The recipe that follows is complete, targets a single 24 GB consumer GPU, and trains all 1.453 B transformer parameters. Its central finding is about which published weights admit it. The model ships in two forms, a rectified-flow checkpoint and the adversarially post-trained checkpoint used for eight-step inference, and a squared-error finetune is valid only on the first. We show that the post-trained weights encode a sample predictor, that a thousand squared-error steps convert them into a conditional-mean predictor, and that the eight-step sampler cannot tolerate the difference. The conversion is invisible to every scalar in the training log — loss falls monotonically throughout — and is measurable in one forward pass by comparing the scale of the one-step estimate to its correlation with the truth. We give that diagnostic, the corrected recipe, and the memory budget that makes it fit.

1 Introduction

Generative audio models have converged on latent diffusion: a high-compression autoencoder maps waveforms to a short sequence of continuous vectors, and a diffusion transformer (DiT) is trained in that space. The recipe is now mature enough that the interesting questions have shifted from how to generate to how to control, and one of the most practically useful controls is continuation: given the first τ seconds of a recording, produce what plausibly follows.

The standard reflex is to treat this as an architec-




Figure 1: The task. A random-length prefix of a latent sequence is retained as context; the model generates the remainder. This is exactly the causal member of the mask family the base model was pretrained with, so no architectural change is required. The training loss is taken only over the generated region.

tural problem. Image inpainting models concatenate a mask and a masked image to the noisy latent, widening the input projection from C to $2C+1$ channels and zero-initialising the new columns. The same construction transfers to audio and works. But it is unnecessary when the base model was already trained with masked conditioning.

Contributions. (i) We identify that outpainting is a subset of the inpainting conditioning already present in Stable Audio 3, and show that restricting the pre-trained mask distribution toward its causal mode is sufficient to specialise the model for continuation (§5.1). (ii) We show that squared-error training on an adversarially post-trained checkpoint silently converts a sample predictor into a conditional-mean predictor, give a one-forward-pass diagnostic that detects it, and report the failure it produces in full (§4, §12). (iii) We argue that the optimiser is part of the recipe rather than an implementation detail, and that an orthogonalised update is what makes a single learning rate valid across a transformer’s matrices (§7). (iv) We give a memory budget under which all 1.453 B transformer parameters, an orthogonalising optimiser and a weight average fit on one 24 GB GPU (§8), and a data path with no manifest and no preprocessing stage (§6). (v) We argue that per-step training loss is the wrong instrument for this objective and propose a frozen evaluation slice (§10) — while showing that even a well-behaved validation curve is not sufficient.

2 Related Work

Latent audio diffusion. Encoding waveforms into a compact latent sequence before diffusion is now standard, and the compression ratio has risen sharply: the autoencoder we build on maps 4096 waveform samples to a single 256-dimensional latent frame, a 10.77 Hz representation of 44.1 kHz stereo. High compression is what makes minute-scale generation tractable, and it is also what makes online encoding expensive (§9).

Masked generative modelling. Inpainting by channel-concatenating a mask and a masked signal originates in image diffusion and has been carried into audio. The variant we exploit differs in that the masked conditioning is injected as a local additive term rather than by widening the input projection, and — critically — is present at pretraining rather than bolted on afterwards.

Flow matching and rectified flow. Rectified flow replaces the variance-preserving diffusion path with a straight interpolation between data and noise, and predicts the constant velocity along it. Its simplicity and its compatibility with few-step distilled samplers have made it the default for recent audio models.

Full versus parameter-efficient finetuning. LoRA and adapters dominate finetuning practice because they fit on small hardware. We take the opposite position for this task: when the target domain differs from the pretraining distribution in timbre and production style rather than in semantics, the low-rank constraint is a real limitation, and a 1.45 B model can be fully finetuned on a consumer card given the right precision and optimiser-state choices.

3 Preliminaries

3.1 Latent representation

Let $\mathbf{x} \in \mathbb{R}^{2 \times N}$ be stereo audio at 44.1 kHz. A frozen autoencoder $\mathcal{E}$ maps it to

$$\mathbf{z} = \mathcal{E}(\mathbf{x}) \in \mathbb{R}^{256 \times T}, \quad T = \lceil N/4096 \rceil, \quad (1)$$

a 256-channel sequence at 10.77 frames per second. The encoder is a patched transformer: waveform samples are first grouped into patches of 256, then a 12-layer transformer with stride 16 produces the latent. The decoder mirrors it. Together they account for 0.852 B parameters and remain frozen throughout.

3.2 Rectified-flow objective

The generative model is a DiT v_θ trained with a rectified-flow (rf_denoiser) objective. For $t \in [0, 1]$ and $\epsilon \sim$




Figure 2: Rectified flow. Data and noise are joined by a straight line, and the network regresses the constant velocity along it. Because the path is linear, the velocity target does not depend on t — but the difficulty of predicting it does, which is the source of the variance discussed in §10.




Figure 3: Timestep density. Mass concentrates in the middle of the schedule where the prediction problem is hardest, and the length-dependent shift moves it toward higher noise as the sequence grows. We inherit both from pretraining rather than substituting a uniform sampler.

$$\mathcal{N}(0, I),$$

$$\mathbf{z}_t = (1 - t) \mathbf{z} + t \epsilon, \quad (2)$$

$$v^* = \epsilon - \mathbf{z}, \quad (3)$$

so that $t=0$ is data and $t=1$ is noise, and the target velocity is constant along each path. The network is conditioned on t , on a text embedding, and on a scalar duration embedding.

3.3 Timestep sampling

Timesteps are drawn from a truncated logistic-normal distribution, reflected, and then warped by a sequence-length-dependent shift σ_L that moves probability mass toward higher noise for longer sequences:

$$t \sim \sigma_L(1 - \text{TLN}(0.075, 1)). \quad (4)$$

We reuse the pretraining sampler unchanged. Matching it matters: a finetune that samples timesteps differently from pretraining spends its gradient budget on a different part of the schedule than the one the model was optimised for.

t	Released		+ 1000 MSE steps	
	std ratio	ρ	std ratio	ρ
0.95	0.832	0.321	0.551	0.502
0.90	0.896	0.565	0.686	0.662
0.80	0.964	0.719	0.784	0.785
0.60	0.999	0.872	0.892	0.899
0.40	1.004	0.946	0.948	0.954
0.20	1.001	0.985	0.983	0.986

Table 1: Scale of the one-step estimate $\hat{\mathbf{z}}_0$ relative to the truth, against its correlation with the truth, on a real excerpt under a full mask. The released checkpoint emits a full-scale sample where it knows almost nothing. After a thousand squared-error steps the two columns coincide at every noise level, which is the defining signature of a conditional-mean predictor.

4 Which checkpoint admits a squared-error finetune

Stability publishes two weights for this model, and the choice between them is the single decision that determines whether the finetune works at all.

Base is the rectified-flow model: it predicts a velocity, and it is sampled with an ODE solver in roughly 50 steps with classifier-free guidance. Released is that model after adversarial post-training. It predicts a sample rather than a conditional mean, and it is sampled with ping-pong in 8 steps at guidance 1:

$$\hat{\mathbf{z}}_0 = \mathbf{z}_t - t v_\theta(\mathbf{z}_t, t, c), \quad (5)$$

$$\mathbf{z}_{t'} = (1 - t') \hat{\mathbf{z}}_0 + t' \mathbf{m}\epsilon', \quad \mathbf{m}\epsilon' \sim \mathcal{N}(0, I). \quad (6)$$

Equation (6) discards the previous iterate entirely and rebuilds it from $\hat{\mathbf{z}}_0$ and fresh noise. Every step therefore requires $\hat{\mathbf{z}}_0$ to be a full-scale, self-consistent clip on its own, at every noise level.

A squared-error objective does not produce that. Its minimiser is $\mathbb{E}[\mathbf{z}_0 \mid \mathbf{z}_t]$, whose scale shrinks exactly in proportion to how much the model actually knows: if ρ is the correlation between the estimate and the truth, then $\text{std}(\hat{\mathbf{z}}_0)/\text{std}(\mathbf{z}_0) \rightarrow \rho$. A sample predictor instead holds that ratio near 1 even where ρ is small. The two are distinguishable in a single forward pass.

Table 1 runs that test. A thousand squared-error steps convert the released checkpoint into a mean predictor: the two columns become equal throughout. At the first ping-pong step the generated signal has lost a third of its amplitude while (6) continues to inject noise at full scale. Perceptually the output flattens and hisses.

Two things make this failure difficult to catch. First, the correlation improves, $0.321 \rightarrow 0.502$ at $t = 0.95$: the model is getting better at the objective it was given, and validation loss falls monotonically throughout. Second, no scalar in the training log distinguishes the two

	$t=0.95$	$t=0.9$	$t=0.8$	$t=0.6$
Base	0.535/0.533	0.605/0.605	0.735/0.740	0.869/0.873
Base + Δ	0.578/0.558	0.639/0.626	0.757/0.753	0.879/0.877
Released	0.920/0.393	0.937/0.483	0.971/0.658	0.999/0.842
Released + Δ	0.965/0.408	0.952/0.506	0.967/0.681	0.995/0.849

Table 2: Amplitude / correlation of the one-step estimate, at 3000 finetuning steps. Equal entries denote a conditional-mean predictor. The transplanted model retains the separation that few-step sampling depends on.

regimes; the diagnostic has to be constructed deliberately.

The mechanism is not incidental — it is the stated purpose of the post-training stage, which works “by supplanting the MSE-based conditional mean loss ... with an adversarial loss”. Running squared error on those weights re-imposes the loss the stage was run to remove, and no learning rate makes that objective correct; a smaller one only slows the regression. Stability’s own trainer refuses the released checkpoint outright, and their documentation gives the remedy: adapters “are trained on the base checkpoint. Once trained, they can be applied to the post-trained model and will work as expected.” We therefore train Base, store checkpoints as the delta from its weights, and add that delta to Released at inference. The same rule is established in images: a distilled model is “impossible to train on directly because every step you train breaks down the compression more and more”.

4.1 The delta transplants

Training Base is only useful if the result can be moved onto Released, and the guidance that says it can is written for low-rank adapters, which are small by construction. A full finetune is not. Fetching one transformer layer of Released by byte range and comparing it with Base puts a number on the margin: adversarial post-training moved that layer’s weights by 1.38% on median, and after 1000 steps our update is already $0.85\times$ that, with two tensors past it. The question is therefore empirical.

Table 2 answers it with the same diagnostic, run on four models under identical noise. Base sits exactly on the conditional-mean line, 0.535 against 0.533, and stays there once our delta is added, which is correct: that is the objective it was trained with. Released separates sharply, and keeps separating after the delta is added: 0.965 against 0.408, against 0.920 and 0.393 for the untouched weights. The few-step behaviour is intact. Its correlation with the truth also rises, so the finetune’s knowledge of the target domain transfers into the accelerated model rather than merely surviving in it.

The practical consequence is that the two checkpoints divide labour cleanly: Base is what the objective is valid




Figure 4: The pretrained mask family. Stable Audio 3 is pretrained with mask-type probabilities $[0.1, 0.8, 0.1]$ over [Spans, Full, Causal]. Outpainting is already in the support of this distribution — it is simply rare. Our finetune reweights to $[0, 0.1, 0.9]$.

against, Released is what inference runs on, and the finetune is transported between them as a float16 delta. Nothing about the recipe needs to change to get both.

5 Method

5.1 Outpainting as a restriction of pretrained masking

The base model consumes two extra conditioning tensors alongside the noisy latent: a binary keep-mask $\mathbf{m} \in \{0, 1\}^{1 \times T}$ and the masked latent $\mathbf{z} \odot \mathbf{m}$. These enter the transformer as a local additive term of dimension $256+1$, added to the token embeddings rather than concatenated to the input channels.

During pretraining, $\mathbf{m}$ is drawn from a family of mask types. Three matter here:

Full $\mathbf{m} = 0$ everywhere. Pure unconditional generation.

Spans a set of non-overlapping interior spans is zeroed. General inpainting.

Causal a prefix length $k \sim \mathcal{U}\{0, \dots, T\}$ is sampled and $\mathbf{m}_i = \mathbf{1}[i < k]$. This is outpainting.

Because Causal is already in the support of the pre-training distribution, the finetune is a reweighting rather than an extension. We set the mask-type probabilities to $[0, 0.1, 0.9]$, concentrating gradient on continuation while retaining a tenth of each batch at Full. That residual matters: training exclusively on causal masks steadily degrades the model’s unconditional behaviour, which is also the behaviour the Full branch of classifier-free guidance relies on at sampling time.

5.2 Two-term objective

The report is explicit that the loss is not taken on the generated region alone: it “is split into two independently averaged terms: a generation loss over the inpainted embeddings ($m = 0$) and a context preservation loss over the kept audio ($m = 1$)”, and the shipped configuration sets the weight on the second to 1.0. Writing $\bar{\mathbf{m}} = 1 - \mathbf{m}$

and $\varepsilon = v_\theta(\mathbf{z}_t, t, c, \mathbf{m}, \mathbf{z} \odot \mathbf{m}) - v^*$,

$$\mathcal{L}(\theta) = \underbrace{\frac{\|\varepsilon \odot \bar{\mathbf{m}}\|_2^2}{256 \sum_i \bar{\mathbf{m}}_i}}_{\text{generation}} + \underbrace{\frac{\|\varepsilon \odot \mathbf{m}\|_2^2}{256 \sum_i \mathbf{m}_i}}_{\text{context}}. \quad (7)$$

Both terms are averaged per sample before the batch mean. Pooling them into a single ratio instead would weight each region by its size, which down-weights precisely the long-context examples that outpainting is about. The context term is cheap — the model is handed the clean latent there, so the target is algebraically recoverable and the term sits near 10^{-2} — but it is what pins the model’s output on the prefix, which the sampler regenerates along with everything else.

5.3 Conditioning is constant, so compute it once

The model is conditioned on a text embedding from a T5Gemma encoder (256 tokens, 768-dimensional) and on a scalar duration embedding. For a single-domain finetune both the prompt and the excerpt length are fixed for the entire run, so the conditioning tensor is a constant. We evaluate the text encoder once at startup, cache the resulting tensor, and release the encoder’s 1.2 GB before the first training step. Classifier-free-guidance dropout is applied downstream by zeroing the cached tensor and therefore needs no encoder.

6 Data path

Conventional latent-diffusion training precomputes latents into a cache. Stability’s own configuration does this (pre_encoded: true). We deliberately do not, and the data path is correspondingly trivial: on each item we open the source file, seek to a uniformly random frame offset, and read exactly the excerpt we need.

This has three consequences worth stating. First, there is no preprocessing stage, no manifest, and no cache to invalidate when the excerpt length or sample rate changes. Second, memory is independent of corpus size — a 26.4-hour, 10.5 GB source file costs nothing to hold because it is never held. Third, the excerpt distribution is continuous rather than quantised to cache-entry boundaries, so the model sees a different alignment of musical structure on every epoch.

The cost is negligible: a random seek-and-decode of a 24s excerpt from a 10.5 GB FLAC takes 31 ms, and two dataloader workers hide it entirely behind GPU work.

7 The optimiser is part of the recipe

The shipped configuration asks for MuonAdamW: Muon on the attention and feed-forward matrices, AdamW on




Figure 5: System overview. Grey blocks are frozen, green is trained. Nothing is precomputed: each step reads fresh audio off disk and encodes it in the same step it is used. The text encoder is evaluated once and then released, since its output is constant for a run.

Algorithm 1 One training step

- 1: $\mathbf{x} \leftarrow \text{RandomExcerpts}(\text{source}, B, \tau)$ $\triangleright$ 31 ms, on workers
 - 2: $\mathbf{z} \leftarrow \mathcal{E}(\mathbf{x})$ $\triangleright$ online, bf16, no grad
 - 3: $\mathbf{m} \leftarrow \text{CausalMask}(B, T; p_{\text{full}})$
 - 4: $t \leftarrow \sigma_L(1 - \text{TLN}())$
 - 5: $\epsilon \sim \mathcal{N}(0, I)$
 - 6: $\mathbf{z}_t \leftarrow (1 - t)\mathbf{z} + t\epsilon; \quad v^* \leftarrow \epsilon - \mathbf{z}$
 - 7: $\hat{v} \leftarrow v_\theta(\mathbf{z}_t, t, c, \mathbf{m}, \mathbf{z} \odot \mathbf{m})$ $\triangleright$ cfg dropout 0.1
 - 8: $\mathcal{L} \leftarrow \|(\hat{v} - v^*) \odot (1 - \mathbf{m})\|^2 / \sum(1 - \mathbf{m})$
 - 9: backward; clip to 1.0; AdamW8bit step
-

everything else. Substituting AdamW throughout is not a neutral simplification.

Muon orthogonalises the momentum buffer before applying it. Writing M for the buffer of a matrix parameter $W \in \mathbb{R}^{m \times n}$ and NS_5 for five steps of the quintic Newton-Schulz iteration,

$$W \leftarrow W - \eta \sqrt{\max(1, m/n)} \text{NS}_5(M + \mu G), \quad (8)$$

so the update has unit spectral norm by construction and η is the size of the step, in the same units, for every matrix in the network. Under AdamW the step size is set by the ratio of the gradient to its own running second moment, which varies by more than an order of magnitude across a transformer’s matrices. A single AdamW learning rate is therefore simultaneously too large for some matrices and too small for others, and the finetune degrades the model along the first group while failing to move it along the second. This is the mechanism behind a finetune whose loss falls while its samples get worse, and it compounds with the checkpoint error of §4.

Fused projections are split before orthogonalisation: a stacked qkv weight is several independent matrices and treating the stack as one is wrong. The configuration’s `fused_layer_patterns` field exists for exactly this. The momentum buffer is held in bfloat16, which halves its cost to 2.9 GB and is harmless because only the direction of (8) survives.

We also restore the exponential moving average the

configuration specifies, $\beta = 0.9995$ under a power-law warmup, and sample from it. It is kept in host memory and updated every eighth step, which costs a copy and returns 2.9 GB of device memory.

8 Fitting a full finetune on one GPU

Training all 1.453 B transformer parameters naively is out of reach: float32 weights, float32 gradients and float32 Adam moments alone require $1.453\text{B} \times 16\text{B} = 23.2\text{GB}$, before activations, before the autoencoder, and before the text encoder. Three choices bring it inside 24 GB.

1. Eight-bit Adam moments. The optimiser state drops from 11.6 GB to 2.9 GB. Master weights and gradients stay in float32, so the update itself is unaffected in precision where it matters.
2. The frozen autoencoder in bfloat16. It never receives gradients, so half precision costs nothing and saves 1.7 GB.
3. Gradient checkpointing throughout the transformer.

Releasing the text encoder after its single use recovers a further 1.2 GB. The resulting measured peak is 15.7 GiB, leaving comfortable headroom for the sampling and decoding done during periodic demos.

8.1 Checkpoint transport

The released checkpoint is 9.2 GB of float32 across 997 tensors (1.453 B transformer, 0.852 B autoencoder). Since training holds float32 master weights in memory regardless, storing bfloat16 on disk costs a single rounding at load. We therefore convert during download, streaming tensors in file order and writing a bfloat16 safetensors file directly, so the float32 artefact is never materialised. The stored model is 4.6 GB.




Figure 6: Memory budget. Eight-bit optimiser moments, a bfloat16 frozen autoencoder and releasing the text encoder move a 27.9 GB configuration to a measured 15.7 GiB peak. Activations under gradient checkpointing are small and not shown separately.

Stage	Time (s)
Excerpt read and decode (2 workers)	0.03
Autoencoder encode (online)	4.00
DiT forward + backward	0.45
Optimiser step	< 0.01
Total per step	4.70

Table 3: Measured per-step cost, RTX 3090, batch 4, 24s excerpts.

9 Cost structure

The autoencoder dominates, costing roughly nine times the transformer step it feeds. This is the honest price of online encoding, and it follows directly from the architecture: a 0.852 B encoder that must process 8085 patch tokens to emit 258 latent frames does substantially more work than a 1.453 B transformer operating on those 258 frames.

We note that this cost is a choice, not a constraint. The training loop consumes $\mathbf{z}$ and is indifferent to its provenance; substituting a precomputed latent cache recovers the 4s and changes nothing else. We prefer the online path for single-corpus finetunes, where setup cost dominates total cost.

10 Evaluation protocol

Rectified-flow training loss at a random timestep is a high-variance estimator of anything one would want to monitor. Figure 8 makes the point concretely: over the first 370 steps of a real run, the per-step loss oscillates between 0.63 and 1.34 and supports no conclusion about whether the model is improving.

We therefore fix an evaluation slice at startup and




Figure 7: Where the step goes. Online encoding is 85% of wall clock. The trade is deliberate: it buys a pipeline with no preprocessing stage and no cache, at roughly 6× the per-step cost of training from precomputed latents.




Figure 8: Training loss is not a learning signal here. The per-step loss (orange) spans 0.63–1.34 with no discernible trend, because each point is a single batch at a single randomly drawn timestep and the difficulty of Eq. (3) varies sharply with t . Holding excerpts, masks, timesteps and noise fixed (blue) yields a monotone curve from the same run.

never resample it: four excerpts drawn from fixed offsets, each with a fixed causal mask at a different prefix fraction (0.25 to 0.75), crossed with four fixed timesteps at $\{0.125, 0.375, 0.625, 0.875\}$, each with fixed noise. The reported value is the mean masked error over these sixteen configurations. Every source of variance except the weights is frozen, so successive values are directly comparable, and the cost is one second per evaluation.

Alongside the scalar we periodically sample a continuation of one fixed reference excerpt with the distilled eight-step ping-pong sampler and log the audio. Fixing the reference is what makes successive clips comparable by ear; sampling a fresh excerpt each time confounds model progress with excerpt difficulty.

11 Recipe summary

Table 5 collects every value needed to reproduce the method. Four entries differ from the base model’s own pretraining configuration: the mask-type probabilities, reweighted toward Causal; the causal context length, drawn in seconds rather than as a fraction of the window; the learning rates, a fifth of the pretraining values

Step	val/loss	Δ	train/loss range
1	1.1005	—	—
50	0.9730	−0.1275	0.98–1.34
100	0.8890	−0.0840	0.87–1.05
150	0.8753	−0.0137	0.85–1.04
200	0.8722	−0.0031	0.83–0.99
250	0.8648	−0.0074	0.86–0.99
300	0.8633	−0.0015	0.63–0.96
350	0.8595	−0.0038	0.86–0.98

Table 4: The frozen slice decreases monotonically while the per-step loss range is wide and overlapping throughout.

Initialisation	Stable Audio 3 Medium Base
Trained	DiT, 1.453 B, all parameters
Frozen	autoencoder 0.852 B, text encoder
Latent	256 ch @ 10.77 Hz (/4096)
Windows	47/95/190/380 s, sampled per step
Batch	1×4 accumulation
Objective	rectified flow, velocity target
Loss	generation + context, Eq. (7)
Mask types	0.55 full / 0.10 segments / 0.35 causal
Causal context	$\mathcal{U}(20, 60)$ seconds
Timesteps	truncated logit-normal, flipped, shifted
CFG dropout	0.1
Optimiser	Muon + AdamW (8-bit moments)
Muon	2×10^{-4} , momentum 0.95, 1.36 B params
AdamW	10^{-5} , (0.9, 0.95), wd 0.01, 94 M params
Schedule	200-step warmup, $\gamma=10^6$, power 0.5
EMA	0.9995, power-law warmup
Grad clip	1.0
Precision	fp32 master, bf16 autocast
Sampling	Euler 50 steps, CFG 4 (Base) ping-pong 8 steps, CFG 1 (Released)
Peak memory	20.2 GiB
Throughput	4.4s/step (RTX 3090)

Table 5: Complete hyperparameter set.

because this is a finetune of a converged model; and eight-bit moments for the AdamW group. Everything else — the objective, the timestep sampler, the distribution shift, the guidance dropout rate, the optimiser pairing, the weight average — is inherited unchanged, which is the point.

12 Diagnosis of a faithful-looking failure

Before arriving at §4 we ran the same recipe against the Released checkpoint, with AdamW in place of the Muon pairing and no weight average. It is worth reporting in full, because every signal available during training said it was working.

We evaluate by outpainting: 32 s of real audio is given as context and the model generates the following 158 s,

Model	Envelope dist.	Level vs. truth
Released, untouched	0.279	1.28×
+ 2000 MSE steps	0.296	1.28×
+ 6000 MSE steps	0.349	0.72×
Blend $\alpha = 0.25$	0.296	1.28×
Blend $\alpha = 0.50$	0.324	1.08×

Table 6: Outcome of squared-error training on the post-trained checkpoint. Quality degrades monotonically with how much of the update is applied, and interpolating back toward the released weights recovers level without ever beating them on envelope.

at the eight sampler steps the released checkpoint is post-trained for. Quality is the mean absolute difference between the generated region’s 24-band log-spectral envelope and that of the true continuation, plus the ratio of generated RMS to the true continuation’s. Five contexts are drawn from widely separated points of a 26.4 h source.

Three ablations exclude the obvious explanations. It is not dosage: stopping at 2000 steps reduces but does not remove the effect. It is not the gates: freezing every 1-D parameter — the RMSNorm gains and adaLN modulation vectors that globally rescale each block — leaves the level ratio at $0.58\times$ against $0.59\times$, so the drift is in the weight matrices, in the velocity field itself. And it is not recoverable by interpolation: a weight blend trades one defect for the other along a curve that never reaches the baseline.

The sampler-step sweep is what points at the mechanism. The released checkpoint degrades gracefully as steps increase ($0.412 \rightarrow 0.577$ from 8 to 64); the finetuned one collapses ($0.221 \rightarrow 2.908$, output falling to 0.017 RMS with spectral centroid rising to 9.2 kHz — hiss). A model whose few-step behaviour is intact but whose many-step behaviour explodes is not a broken model; it is a model being sampled by the wrong solver family. Table 1 then measures the cause directly.

Throughout all of this, validation loss fell monotonically, from 1.31 to 1.01, and the context-reconstruction term fell twenty-fold. The training curves were indistinguishable from a healthy run.

13 Discussion

The headline lesson is worth stating plainly: a finetune can be faithful to its reference implementation, train stably, and lower its loss while making the model worse, and the loss curve cannot tell you. Ours fell from 1.31 to 1.01 over exactly the interval in which the model was being converted from a sample predictor into a mean predictor. Any evaluation that had stopped at the scalar would have reported success.

The generalisable part is the diagnostic, not the in-

cident. Whenever a sampler rebuilds its state from a one-step estimate — ping-pong, restart sampling, consistency-model sampling — that estimate must be a sample and not a mean, and the distinction is measurable in one forward pass by comparing the estimate’s scale to its correlation with the truth. The two coincide for a mean predictor and separate for a sample predictor. We would run this before and after any training that touches a distilled or post-trained checkpoint.

The general lesson is that capability is often already latent in a pretrained checkpoint, distributed across a conditioning interface that was trained for a broader task than the one at hand. Stable Audio 3 was trained to inpaint; outpainting is the causal corner of that distribution. Specialising to it is a change of sampling weights, not a change of model. The practical payoff is disproportionate: no zero-initialised branch to warm up, no train/inference conditioning mismatch, and a finetune whose step zero is exactly the released model.

Two limitations are worth naming. Online encoding costs 85% of wall-clock time, which is the right trade for single-corpus finetunes and the wrong one for large multi-corpus runs. And a full finetune on one domain will move the model away from its general prior; the residual Full-mask fraction slows this but does not eliminate it.

14 Conclusion

We presented a recipe for specialising a latent rectified-flow audio diffusion model to outpainting, with no architectural modification, on a single consumer GPU. The method reduces to reweighting a pretrained mask distribution, keeping the reference’s two-term objective and its orthogonalising optimiser, and choosing precisions that let 1.453B parameters train in 24GB.

The result we would most want carried forward is the negative one and its instrument. Squared-error training on an adversarially post-trained checkpoint is not merely suboptimal; it systematically undoes the post-training, and it does so while every scalar available during the run improves. The distinction between a sample predictor and a conditional-mean predictor is not visible in a loss curve, but it is visible in one forward pass, and it should be checked before and after any training that touches a distilled checkpoint. Code is available at <https://github.com/mrcolo/stems-loop>.